

Assignment-12: TensorFlow and Keras: Build various MLP architectures for MNIST dataset

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
/home/ubuntu/anaconda3/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:516:
FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future
version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:517:
FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future
version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:518:
FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future
version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:519:
FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future
version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_quint16 = np.dtype [("quint16", np.uint16, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:520:
FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future
version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_qint32 = np.dtype [("qint32", np.int32, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:525:
FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future
version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
_np_resource = np.dtype [("resource", np.ubyte, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-
packages/tensorboard/compat/tensorflow_stub/dtypes.py:541: FutureWarning: Passing (type, 1) or 'lt
ype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (t
ype, (1,)) / '(1,)type'.
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-
packages/tensorboard/compat/tensorflow_stub/dtypes.py:542: FutureWarning: Passing (type, 1) or 'lt
ype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (t
ype, (1,)) / '(1,)type'.
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-
packages/tensorboard/compat/tensorflow_stub/dtypes.py:543: FutureWarning: Passing (type, 1) or 'lt
ype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (t
ype, (1,)) / '(1,)type'.
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-
packages/tensorboard/compat/tensorflow_stub/dtypes.py:544: FutureWarning: Passing (type, 1) or 'lt
ype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (t
ype, (1,)) / '(1,)type'.
_np_quint16 = np.dtype [("quint16", np.uint16, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-
packages/tensorboard/compat/tensorflow_stub/dtypes.py:545: FutureWarning: Passing (type, 1) or 'lt
ype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (t
ype, (1,)) / '(1,)type'.
_np_qint32 = np.dtype [("qint32", np.int32, 1)]
/home/ubuntu/anaconda3/lib/python3.6/site-
packages/tensorboard/compat/tensorflow_stub/dtypes.py:550: FutureWarning: Passing (type, 1) or 'lt
ype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (t
ype, (1,)) / '(1,)type'.
_np_resource = np.dtype [("resource", np.ubyte, 1)]
```

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
```

```

import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()

```

In [3]:

```

# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()

```

In [4]:

```

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))

```

Number of training examples : 60000 and each image is of shape (28, 28)
 Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```

# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])

```

In [6]:

```

# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))

```

Number of training examples : 60000 and each image is of shape (784)
 Number of training examples : 10000 and each image is of shape (784)

In [7]:

```

# An example data point
print(X_train[0])

```

```

[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205 11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0  0

```


0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686

In [10]:

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

In [11]:

Dense implements the operation: $\text{output} = \text{activation}(\text{dot}(\text{input}, \text{kernel}) + \text{bias})$ where


```

WARNING: Logging before flag parsing goes to stderr.
W0818 11:16:39.769354 140217894131520 deprecation_wrapper.py:119] From
/home/ubuntu/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:74: The nam
e tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

W0818 11:16:39.781115 140217894131520 deprecation_wrapper.py:119] From
/home/ubuntu/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:517: The na
me tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

W0818 11:16:39.783540 140217894131520 deprecation_wrapper.py:119] From
/home/ubuntu/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:4185: The n
ame tf.truncated_normal is deprecated. Please use tf.random.truncated_normal instead.

W0818 11:16:39.800450 140217894131520 deprecation_wrapper.py:119] From
/home/ubuntu/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:4138: The n
ame tf.random_uniform is deprecated. Please use tf.random.uniform instead.

W0818 11:16:39.810351 140217894131520 deprecation_wrapper.py:119] From
/home/ubuntu/anaconda3/lib/python3.6/site-packages/keras/optimizers.py:790: The name
tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

W0818 11:16:39.827991 140217894131520 deprecation_wrapper.py:119] From
/home/ubuntu/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3295: The n
ame tf.log is deprecated. Please use tf.math.log instead.

W0818 11:16:39.893717 140217894131520 deprecation.py:323] From
/home/ubuntu/anaconda3/lib/python3.6/site-packages/tensorflow/python/ops/math_grad.py:1250:
add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where

```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 472)	370520
dense_2 (Dense)	(None, 168)	79464
dense_3 (Dense)	(None, 10)	1690
Total params: 451,674		
Trainable params: 451,674		
Non-trainable params: 0		

```

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 36us/step - loss: 0.2344 - acc: 0.9320 -
val_loss: 0.1033 - val_acc: 0.9694
Epoch 2/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0872 - acc: 0.9738 -
val_loss: 0.0851 - val_acc: 0.9757
Epoch 3/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0558 - acc: 0.9827 -
val_loss: 0.0665 - val_acc: 0.9795
Epoch 4/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0365 - acc: 0.9881 -
val_loss: 0.0719 - val_acc: 0.9770
Epoch 5/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0281 - acc: 0.9911 -
val_loss: 0.0605 - val_acc: 0.9811
Epoch 6/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0214 - acc: 0.9930 -
val_loss: 0.0674 - val_acc: 0.9808
Epoch 7/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0178 - acc: 0.9940 -
val_loss: 0.0681 - val_acc: 0.9809
Epoch 8/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0135 - acc: 0.9955 -
val_loss: 0.0706 - val_acc: 0.9802
Epoch 9/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0122 - acc: 0.9960 -
val_loss: 0.0721 - val_acc: 0.9819
Epoch 10/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0133 - acc: 0.9959 -

```

```

00000/00000 [-----] - 1s 18us/step - loss: 0.0133 - acc: 0.9959 -
val_loss: 0.0800 - val_acc: 0.9800
Epoch 11/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0110 - acc: 0.9963 -
val_loss: 0.0796 - val_acc: 0.9819
Epoch 12/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0084 - acc: 0.9971 -
val_loss: 0.0969 - val_acc: 0.9799
Epoch 13/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0091 - acc: 0.9968 -
val_loss: 0.0929 - val_acc: 0.9820
Epoch 14/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0086 - acc: 0.9970 -
val_loss: 0.0888 - val_acc: 0.9822
Epoch 15/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0106 - acc: 0.9965 -
val_loss: 0.0982 - val_acc: 0.9788
Epoch 16/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0062 - acc: 0.9980 -
val_loss: 0.0976 - val_acc: 0.9789
Epoch 17/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0062 - acc: 0.9981 -
val_loss: 0.0926 - val_acc: 0.9804
Epoch 18/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0088 - acc: 0.9971 -
val_loss: 0.0860 - val_acc: 0.9814
Epoch 19/20
60000/60000 [=====] - 1s 20us/step - loss: 0.0047 - acc: 0.9986 -
val_loss: 0.0883 - val_acc: 0.9816
Epoch 20/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0088 - acc: 0.9971 -
val_loss: 0.1011 - val_acc: 0.9790

```

In [15]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
score1=score[0]
score2=score[1]
train_acc1=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax11 = plt.subplots(1,1)
ax11.set_xlabel('epoch') ; ax11.set_ylabel('Categorical Cross-entropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(x_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(x_test, y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy11 = history11.history['val_loss']
ty11 = history11.history['loss']
plt_dynamic(x, vy11, ty11, ax11)

```

Test score: 0.10111289650749382

Test accuracy: 0.979





In [16]:


```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





MLP + ReLU activation function + RMSprop optimizer

In [17]:

```

from keras.initializers import he_normal
import warnings
warnings.filterwarnings("ignore")
model_relu = Sequential()
model_relu.add(Dense(472, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_relu.add(Dense(168, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history11 = model_relu.fit(X_train, Y_train,
                          batch_size=batch_size,
                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 472)	370520
dense_5 (Dense)	(None, 168)	79464
dense_6 (Dense)	(None, 10)	1690
Total params: 451,674		
Trainable params: 451,674		
Non-trainable params: 0		

```

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 1s 20us/step - loss: 0.2273 - acc: 0.9312 -
val_loss: 0.1277 - val_acc: 0.9582
Epoch 2/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0854 - acc: 0.9732 -
val_loss: 0.0855 - val_acc: 0.9714
Epoch 3/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0568 - acc: 0.9825 -
val_loss: 0.0737 - val_acc: 0.9769
Epoch 4/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0405 - acc: 0.9873 -
val_loss: 0.0619 - val_acc: 0.9794
Epoch 5/20
60000/60000 [=====] - 1s 17us/step - loss: 0.0295 - acc: 0.9906 -
val_loss: 0.0759 - val_acc: 0.9796
Epoch 6/20
60000/60000 [=====] - 1s 20us/step - loss: 0.0230 - acc: 0.9928 -
val_loss: 0.0871 - val_acc: 0.9763
Epoch 7/20
60000/60000 [=====] - 1s 20us/step - loss: 0.0174 - acc: 0.9942 -
val_loss: 0.0786 - val_acc: 0.9821
Epoch 8/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0146 - acc: 0.9959 -
val_loss: 0.0832 - val_acc: 0.9823
Epoch 9/20
60000/60000 [=====] - 1s 17us/step - loss: 0.0121 - acc: 0.9963 -
val_loss: 0.0794 - val_acc: 0.9830
Epoch 10/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0096 - acc: 0.9969 -

```

```

val_loss: 0.0941 - val_acc: 0.9788
Epoch 11/20
60000/60000 [=====] - 1s 20us/step - loss: 0.0083 - acc: 0.9973 -
val_loss: 0.1075 - val_acc: 0.9791
Epoch 12/20
60000/60000 [=====] - 1s 20us/step - loss: 0.0082 - acc: 0.9976 -
val_loss: 0.0988 - val_acc: 0.9817
Epoch 13/20
60000/60000 [=====] - 1s 20us/step - loss: 0.0061 - acc: 0.9981 -
val_loss: 0.0931 - val_acc: 0.9820
Epoch 14/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0060 - acc: 0.9983 -
val_loss: 0.1080 - val_acc: 0.9810
Epoch 15/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0053 - acc: 0.9984 -
val_loss: 0.1159 - val_acc: 0.9807
Epoch 16/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0042 - acc: 0.9988 -
val_loss: 0.1157 - val_acc: 0.9817
Epoch 17/20
60000/60000 [=====] - 1s 19us/step - loss: 0.0048 - acc: 0.9986 -
val_loss: 0.1382 - val_acc: 0.9806
Epoch 18/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0050 - acc: 0.9983 -
val_loss: 0.1184 - val_acc: 0.9832
Epoch 19/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0038 - acc: 0.9989 -
val_loss: 0.1224 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 1s 18us/step - loss: 0.0031 - acc: 0.9991 -
val_loss: 0.1398 - val_acc: 0.9809

```

In [18]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
score1=score[0]
score2=score[1]
train_accl=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax11 = plt.subplots(1,1)
ax11.set_xlabel('epoch') ; ax11.set_ylabel('Categorical Cross-entropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(x_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(x_test, y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy11 = history11.history['val_loss']
ty11 = history11.history['loss']
plt_dynamic(x, vy11, ty11, ax11)

```

Test score: 0.13976741824084157
Test accuracy: 0.9809





In [19]:


```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





1.2 MLP + Batch-Norm on hidden Layers + Adam Optimizer

In [20]:

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(472, activation='relu',
                      input_shape=(input_dim,),
                      kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(168, activation='relu',
                      kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 472)	370520
batch_normalization_1 (Batch Normalization)	(None, 472)	1888
dense_8 (Dense)	(None, 168)	79464
batch_normalization_2 (Batch Normalization)	(None, 168)	672
dense_9 (Dense)	(None, 10)	1690
Total params: 454,234		
Trainable params: 452,954		
Non-trainable params: 1,280		

In [21]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy',
                   metrics=['accuracy'])

history12 = model_batch.fit(X_train, Y_train,
                           batch_size=batch_size,
                           epochs=nb_epoch, verbose=1,
                           validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 41us/step - loss: 0.1872 - acc: 0.9437 -
val_loss: 0.1048 - val_acc: 0.9668
Epoch 2/20
60000/60000 [=====] - 2s 30us/step - loss: 0.0699 - acc: 0.9789 -
val_loss: 0.0926 - val_acc: 0.9712
Epoch 3/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0446 - acc: 0.9866 -
val_loss: 0.0824 - val_acc: 0.9739
Epoch 4/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0322 - acc: 0.9902 -
val_loss: 0.0861 - val_acc: 0.9754
Epoch 5/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0273 - acc: 0.9913 -
val_loss: 0.0790 - val_acc: 0.9755
Epoch 6/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0211 - acc: 0.9930 -
```

```

00000/00000 [-----] - 2s 31us/step - loss: 0.0211 - acc: 0.9939
val_loss: 0.0764 - val_acc: 0.9787
Epoch 7/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0197 - acc: 0.9933 -
val_loss: 0.0752 - val_acc: 0.9793
Epoch 8/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0175 - acc: 0.9944 -
val_loss: 0.0679 - val_acc: 0.9802
Epoch 9/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0134 - acc: 0.9955 -
val_loss: 0.0710 - val_acc: 0.9818
Epoch 10/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0132 - acc: 0.9960 -
val_loss: 0.0781 - val_acc: 0.9791
Epoch 11/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0143 - acc: 0.9952 -
val_loss: 0.0780 - val_acc: 0.9774
Epoch 12/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0111 - acc: 0.9963 -
val_loss: 0.0877 - val_acc: 0.9781
Epoch 13/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0101 - acc: 0.9968 -
val_loss: 0.0709 - val_acc: 0.9817
Epoch 14/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0074 - acc: 0.9977 -
val_loss: 0.0819 - val_acc: 0.9796
Epoch 15/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0118 - acc: 0.9960 -
val_loss: 0.0832 - val_acc: 0.9799
Epoch 16/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0099 - acc: 0.9965 -
val_loss: 0.0708 - val_acc: 0.9816
Epoch 17/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0089 - acc: 0.9968 -
val_loss: 0.0942 - val_acc: 0.9792
Epoch 18/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0088 - acc: 0.9971 -
val_loss: 0.0756 - val_acc: 0.9807
Epoch 19/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0055 - acc: 0.9982 -
val_loss: 0.0713 - val_acc: 0.9829
Epoch 20/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0056 - acc: 0.9982 -
val_loss: 0.0893 - val_acc: 0.9797

```

In [22]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
score3=score[0]
score4=score[1]
train_acc2=history11.history['acc']

fig,ax12 = plt.subplots(1,1)
ax12.set_xlabel('epoch') ; ax12.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


vy12 = history12.history['val_loss']
ty12 = history12.history['loss']
plt_dynamic(x, vy12, ty12, ax12)

```

Test score: 0.08932151498529384

Test accuracy: 0.9797





In [23]:


```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hidden Layer 1

Hidden Layer 2

Output Layer

MLP + Batch-Norm on hidden Layers + Adagrad Optimizer

In [24]:

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(472, activation='relu',
                      input_shape=(input_dim,),
                      kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(168, activation='relu',
                      kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.compile(optimizer='adagrad', loss='categorical_crossentropy',
                    metrics=['accuracy'])

history12 = model_batch.fit(X_train, Y_train,
                             batch_size=batch_size,
                             epochs=nb_epoch, verbose=1,
                             validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 2s 36us/step - loss: 0.1549 - acc: 0.9541 -
val_loss: 0.0866 - val_acc: 0.9746
Epoch 2/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0572 - acc: 0.9836 -
val_loss: 0.0676 - val_acc: 0.9795
Epoch 3/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0339 - acc: 0.9912 -
val_loss: 0.0652 - val_acc: 0.9803
Epoch 4/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0220 - acc: 0.9950 -
val_loss: 0.0593 - val_acc: 0.9823
Epoch 5/20
60000/60000 [=====] - 2s 30us/step - loss: 0.0142 - acc: 0.9976 -
val_loss: 0.0601 - val_acc: 0.9813
Epoch 6/20
60000/60000 [=====] - 2s 30us/step - loss: 0.0101 - acc: 0.9986 -
val_loss: 0.0592 - val_acc: 0.9826
Epoch 7/20
60000/60000 [=====] - 2s 30us/step - loss: 0.0074 - acc: 0.9992 -
val_loss: 0.0571 - val_acc: 0.9833
Epoch 8/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0057 - acc: 0.9994 -
val_loss: 0.0589 - val_acc: 0.9829
Epoch 9/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0045 - acc: 0.9998 -
val_loss: 0.0589 - val_acc: 0.9831
Epoch 10/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0038 - acc: 0.9999 -
val_loss: 0.0584 - val_acc: 0.9836
Epoch 11/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0033 - acc: 0.9999 -
val_loss: 0.0592 - val_acc: 0.9833
Epoch 12/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0028 - acc: 0.9999 -
val_loss: 0.0576 - val_acc: 0.9831
Epoch 13/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0024 - acc: 1.0000 -
val_loss: 0.0586 - val_acc: 0.9825
Epoch 14/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0021 - acc: 1.0000 -
val_loss: 0.0578 - val_acc: 0.9836
```



```

Epoch 15/20
60000/60000 [=====] - 2s 29us/step - loss: 0.0019 - acc: 1.0000 -
val_loss: 0.0588 - val_acc: 0.9833
Epoch 16/20
60000/60000 [=====] - 2s 29us/step - loss: 0.0017 - acc: 1.0000 -
val_loss: 0.0588 - val_acc: 0.9842
Epoch 17/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0015 - acc: 1.0000 -
val_loss: 0.0601 - val_acc: 0.9841
Epoch 18/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0014 - acc: 1.0000 -
val_loss: 0.0599 - val_acc: 0.9844
Epoch 19/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0014 - acc: 1.0000 -
val_loss: 0.0607 - val_acc: 0.9836
Epoch 20/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0012 - acc: 1.0000 -
val_loss: 0.0604 - val_acc: 0.9843

```

In [25]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
score3=score[0]
score4=score[1]
train_acc2=history11.history['acc']

fig,ax12 = plt.subplots(1,1)
ax12.set_xlabel('epoch') ; ax12.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy12 = history12.history['val_loss']
ty12 = history12.history['loss']
plt_dynamic(x, vy12, ty12, ax12)

```

Test score: 0.060404907803970856

Test accuracy: 0.9843



In [26]:

```

w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



1.3 MLP + Dropout + AdamOptimizer

In [27]:

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(472, activation='relu',
                    input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(168, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )

```

```

model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

W0818 11:21:10.975695 140217894131520 deprecation.py:506] From /home/ubuntu/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 472)	370520
batch_normalization_5 (Batch Normalization)	(None, 472)	1888
dropout_1 (Dropout)	(None, 472)	0
dense_14 (Dense)	(None, 168)	79464
batch_normalization_6 (Batch Normalization)	(None, 168)	672
dropout_2 (Dropout)	(None, 168)	0
dense_15 (Dense)	(None, 10)	1690
Total params: 454,234		
Trainable params: 452,954		
Non-trainable params: 1,280		

In [28]:

```

model_drop.compile(optimizer='adam',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])

history13 = model_drop.fit(X_train, Y_train,
                           batch_size=batch_size,

                           epochs=nb_epoch, verbose=1,
                           validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 3s 46us/step - loss: 0.4263 - acc: 0.8711 - val_loss: 0.1414 - val_acc: 0.9557

Epoch 2/20
60000/60000 [=====] - 2s 32us/step - loss: 0.2022 - acc: 0.9400 - val_loss: 0.1110 - val_acc: 0.9667

Epoch 3/20
60000/60000 [=====] - 2s 32us/step - loss: 0.1594 - acc: 0.9509 - val_loss: 0.0941 - val_acc: 0.9705

Epoch 4/20
60000/60000 [=====] - 2s 36us/step - loss: 0.1349 - acc: 0.9581 - val_loss: 0.0790 - val_acc: 0.9745

Epoch 5/20
60000/60000 [=====] - 2s 36us/step - loss: 0.1176 - acc: 0.9637 - val_loss: 0.0763 - val_acc: 0.9759

Epoch 6/20
60000/60000 [=====] - 2s 35us/step - loss: 0.1087 - acc: 0.9663 - val_loss: 0.0764 - val_acc: 0.9776

Epoch 7/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0976 - acc: 0.9701 - val_loss: 0.0635 - val_acc: 0.9806

Epoch 8/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0899 - acc: 0.9726 - val_loss: 0.0679 - val_acc: 0.9797

```

Epoch 9/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0877 - acc: 0.9728 -
val_loss: 0.0610 - val_acc: 0.9810
Epoch 10/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0817 - acc: 0.9742 -
val_loss: 0.0598 - val_acc: 0.9823
Epoch 11/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0803 - acc: 0.9746 -
val_loss: 0.0587 - val_acc: 0.9812
Epoch 12/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0730 - acc: 0.9769 -
val_loss: 0.0600 - val_acc: 0.9811
Epoch 13/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0723 - acc: 0.9773 -
val_loss: 0.0641 - val_acc: 0.9807
Epoch 14/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0690 - acc: 0.9780 -
val_loss: 0.0579 - val_acc: 0.9827
Epoch 15/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0631 - acc: 0.9801 -
val_loss: 0.0569 - val_acc: 0.9828
Epoch 16/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0598 - acc: 0.9808 -
val_loss: 0.0532 - val_acc: 0.9831
Epoch 17/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0576 - acc: 0.9810 -
val_loss: 0.0576 - val_acc: 0.9830
Epoch 18/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0596 - acc: 0.9811 -
val_loss: 0.0573 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0534 - acc: 0.9830 -
val_loss: 0.0641 - val_acc: 0.9821
Epoch 20/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0515 - acc: 0.9834 -
val_loss: 0.0559 - val_acc: 0.9836

```

In [29]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
score5=score[0]
score6=score[1]
train_acc3=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax13 = plt.subplots(1,1)
ax13.set_xlabel('epoch') ; ax13.set_ylabel('Categorical Crossentropy Loss')


vy13 = history13.history['val_loss']
ty13 = history13.history['loss']
plt_dynamic(x, vy13, ty13, ax13)

```

Test score: 0.05593528219778964

Test accuracy: 0.9836





In [30]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout + Adadelata Optimizer

In [31]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

keras

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(472, activation='relu',
                    input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(168, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.4))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 472)	370520
batch_normalization_7 (Batch Normalization)	(None, 472)	1888
dropout_3 (Dropout)	(None, 472)	0
dense_17 (Dense)	(None, 168)	79464
batch_normalization_8 (Batch Normalization)	(None, 168)	672
dropout_4 (Dropout)	(None, 168)	0
dense_18 (Dense)	(None, 10)	1690
Total params: 454,234		
Trainable params: 452,954		
Non-trainable params: 1,280		

In [32]:

```
model_drop.compile(optimizer='adadelta',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history13 = model_drop.fit(X_train, Y_train,
                          batch_size=batch_size,

                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 3s 47us/step - loss: 0.3931 - acc: 0.8808 -
val_loss: 0.1345 - val_acc: 0.9589
Epoch 2/20
60000/60000 [=====] - 2s 35us/step - loss: 0.1973 - acc: 0.9409 -
val_loss: 0.1051 - val_acc: 0.9677
Epoch 3/20
60000/60000 [=====] - 2s 34us/step - loss: 0.1523 - acc: 0.9533 -
val_loss: 0.0853 - val_acc: 0.9736
Epoch 4/20
60000/60000 [=====] - 2s 34us/step - loss: 0.1245 - acc: 0.9615 -
val_loss: 0.0800 - val_acc: 0.9757
Epoch 5/20
60000/60000 [=====] - 2s 34us/step - loss: 0.1114 - acc: 0.9657 -
val_loss: 0.0712 - val_acc: 0.9776
Epoch 6/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0985 - acc: 0.9698 -
val_loss: 0.0675 - val_acc: 0.9793
```

```

Epoch 7/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0880 - acc: 0.9724 -
val_loss: 0.0633 - val_acc: 0.9805
Epoch 8/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0825 - acc: 0.9747 -
val_loss: 0.0631 - val_acc: 0.9807
Epoch 9/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0752 - acc: 0.9770 -
val_loss: 0.0594 - val_acc: 0.9816
Epoch 10/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0711 - acc: 0.9779 -
val_loss: 0.0560 - val_acc: 0.9828
Epoch 11/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0675 - acc: 0.9789 -
val_loss: 0.0619 - val_acc: 0.9810
Epoch 12/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0636 - acc: 0.9798 -
val_loss: 0.0556 - val_acc: 0.9831
Epoch 13/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0616 - acc: 0.9802 -
val_loss: 0.0560 - val_acc: 0.9831
Epoch 14/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0537 - acc: 0.9829 -
val_loss: 0.0559 - val_acc: 0.9834
Epoch 15/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0549 - acc: 0.9825 -
val_loss: 0.0562 - val_acc: 0.9832
Epoch 16/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0507 - acc: 0.9838 -
val_loss: 0.0563 - val_acc: 0.9833
Epoch 17/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0489 - acc: 0.9847 -
val_loss: 0.0542 - val_acc: 0.9838
Epoch 18/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0465 - acc: 0.9850 -
val_loss: 0.0542 - val_acc: 0.9839
Epoch 19/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0453 - acc: 0.9853 -
val_loss: 0.0539 - val_acc: 0.9847
Epoch 20/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0451 - acc: 0.9849 -
val_loss: 0.0525 - val_acc: 0.9840

```

In [33]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
score5=score[0]
score6=score[1]
train_acc3=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax13 = plt.subplots(1,1)
ax13.set_xlabel('epoch') ; ax13.set_ylabel('Categorical Crossentropy Loss')


vy13 = history13.history['val_loss']
ty13 = history13.history['loss']
plt_dynamic(x, vy13, ty13, ax13)

```

Test score: 0.05253026305373642

Test accuracy: 0.984





In [34]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2) 3-Hidden layer architecture (784-352-164-124 architecture)

2.1 MLP + ReLU + ADAM

In [35]:

```
model_relu = Sequential()
model_relu.add(Dense(352, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_relu.add(Dense(164, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )

model_relu.add(Dense(124, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history21 = model_relu.fit(X_train, Y_train,
                          batch_size=batch_size,
                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_19 (Dense)	(None, 352)	276320
dense_20 (Dense)	(None, 164)	57892
dense_21 (Dense)	(None, 124)	20460
dense_22 (Dense)	(None, 10)	1250

=====
Total params: 355,922
Trainable params: 355,922
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 2s 33us/step - loss: 0.2438 - acc: 0.9264 - val_loss: 0.1182 - val_acc: 0.9642

Epoch 2/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0875 - acc: 0.9728 - val_loss: 0.0852 - val_acc: 0.9727

Epoch 3/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0578 - acc: 0.9819 - val_loss: 0.0889 - val_acc: 0.9729

Epoch 4/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0429 - acc: 0.9865 - val_loss: 0.0751 - val_acc: 0.9788

Epoch 5/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0309 - acc: 0.9896 - val_loss: 0.0666 - val_acc: 0.9815

Epoch 6/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0241 - acc: 0.9918 - val_loss: 0.0705 - val_acc: 0.9808

Epoch 7/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0204 - acc: 0.9932 - val_loss: 0.1083 - val_acc: 0.9745

Epoch 8/20
60000/60000 [=====] - 1s 23us/step - loss: 0.0199 - acc: 0.9933 - val_loss: 0.0818 - val_acc: 0.9803

Epoch 9/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0168 - acc: 0.9944 - val_loss: 0.0873 - val_acc: 0.9789

Epoch 10/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0152 - acc: 0.9948 - val_loss: 0.0907 - val_acc: 0.9777

```

Epoch 11/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0136 - acc: 0.9957 -
val_loss: 0.0779 - val_acc: 0.9815
Epoch 12/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0131 - acc: 0.9956 -
val_loss: 0.1167 - val_acc: 0.9751
Epoch 13/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0121 - acc: 0.9959 -
val_loss: 0.1140 - val_acc: 0.9779
Epoch 14/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0123 - acc: 0.9957 -
val_loss: 0.0964 - val_acc: 0.9797
Epoch 15/20
60000/60000 [=====] - 1s 21us/step - loss: 0.0100 - acc: 0.9968 -
val_loss: 0.1011 - val_acc: 0.9794
Epoch 16/20
60000/60000 [=====] - 1s 23us/step - loss: 0.0117 - acc: 0.9960 -
val_loss: 0.1040 - val_acc: 0.9803
Epoch 17/20
60000/60000 [=====] - 1s 25us/step - loss: 0.0108 - acc: 0.9965 -
val_loss: 0.1029 - val_acc: 0.9791
Epoch 18/20
60000/60000 [=====] - 1s 23us/step - loss: 0.0066 - acc: 0.9979 -
val_loss: 0.0977 - val_acc: 0.9798
Epoch 19/20
60000/60000 [=====] - 1s 23us/step - loss: 0.0097 - acc: 0.9969 -
val_loss: 0.0929 - val_acc: 0.9825
Epoch 20/20
60000/60000 [=====] - 1s 22us/step - loss: 0.0071 - acc: 0.9977 -
val_loss: 0.1009 - val_acc: 0.9798

```

In [36]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
score7=score[0]
score8=score[1]
train_acc4=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax21 = plt.subplots(1,1)
ax21.set_xlabel('epoch') ; ax21.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))


vy21 = history21.history['val_loss']
ty21 = history21.history['loss']
plt_dynamic(x, vy21, ty21, ax21)

```

Test score: 0.10094069364849065

Test accuracy: 0.9798





In [38]:

```
w_after = model_relu.get_weights()

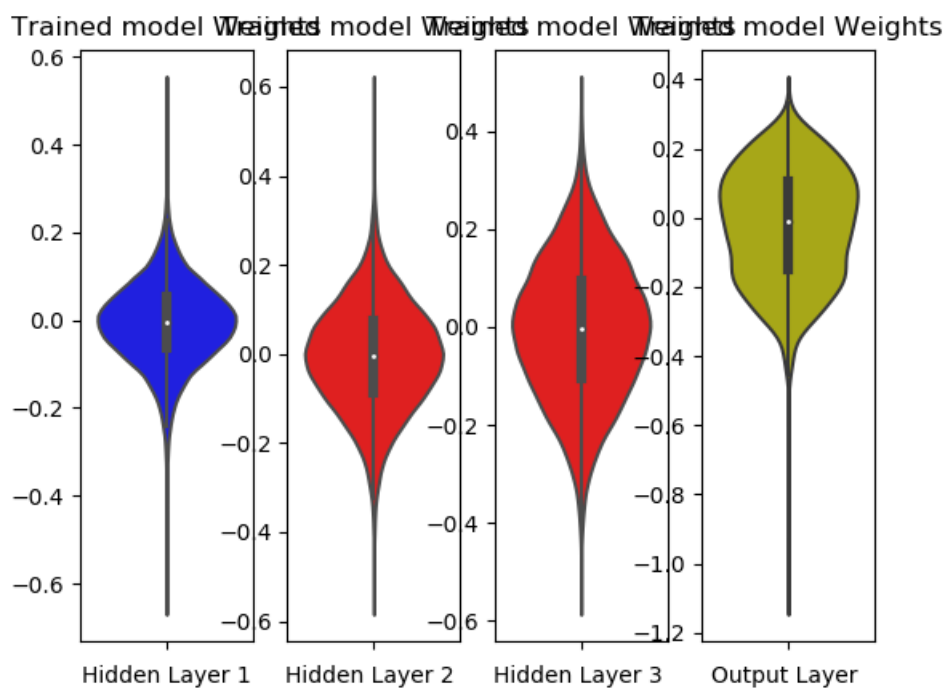
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU + Adadelta

In [39]:

```
model_relu = Sequential()
model_relu.add(Dense(352, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_relu.add(Dense(164, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )

model_relu.add(Dense(124, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adadelta',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history21 = model_relu.fit(X_train, Y_train,
                          batch_size=batch_size,
                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_23 (Dense)	(None, 352)	276320
dense_24 (Dense)	(None, 164)	57892
dense_25 (Dense)	(None, 124)	20460
dense_26 (Dense)	(None, 10)	1250

=====
Total params: 355,922
Trainable params: 355,922
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 34us/step - loss: 0.2696 - acc: 0.9177 -

val_loss: 0.1526 - val_acc: 0.9511

Epoch 2/20

60000/60000 [=====] - 1s 22us/step - loss: 0.1003 - acc: 0.9694 -

val_loss: 0.1111 - val_acc: 0.9633

Epoch 3/20

60000/60000 [=====] - 1s 22us/step - loss: 0.0657 - acc: 0.9804 -

val_loss: 0.0967 - val_acc: 0.9707

Epoch 4/20

60000/60000 [=====] - 1s 22us/step - loss: 0.0460 - acc: 0.9857 -

val_loss: 0.0750 - val_acc: 0.9766

Epoch 5/20

60000/60000 [=====] - 1s 22us/step - loss: 0.0325 - acc: 0.9897 -

val_loss: 0.0682 - val_acc: 0.9802

Epoch 6/20

60000/60000 [=====] - 1s 22us/step - loss: 0.0234 - acc: 0.9927 -

val_loss: 0.0776 - val_acc: 0.9785

Epoch 7/20

60000/60000 [=====] - 1s 22us/step - loss: 0.0169 - acc: 0.9949 -

val_loss: 0.0612 - val_acc: 0.9828

Epoch 8/20

60000/60000 [=====] - 2s 25us/step - loss: 0.0119 - acc: 0.9964 -

val_loss: 0.0760 - val_acc: 0.9798

Epoch 9/20

60000/60000 [=====] - 1s 22us/step - loss: 0.0082 - acc: 0.9977 -

val_loss: 0.0690 - val_acc: 0.9829

Epoch 10/20

60000/60000 [=====] - 1s 22us/step - loss: 0.0062 - acc: 0.9983 -

val_loss: 0.0698 - val_acc: 0.9841

Epoch 11/20

60000/60000 [=====] - 1s 22us/step - loss: 0.0043 - acc: 0.9988 -

val_loss: 0.0805 - val_acc: 0.9817

Epoch 12/20

```

60000/60000 [=====] - 1s 22us/step - loss: 0.0021 - acc: 0.9995 -
val_loss: 0.0769 - val_acc: 0.9828
Epoch 13/20
60000/60000 [=====] - 1s 22us/step - loss: 0.0017 - acc: 0.9996 -
val_loss: 0.0793 - val_acc: 0.9824
Epoch 14/20
60000/60000 [=====] - 1s 24us/step - loss: 5.2830e-04 - acc: 0.9999 - val
_loss: 0.0790 - val_acc: 0.9837
Epoch 15/20
60000/60000 [=====] - 1s 22us/step - loss: 2.4163e-04 - acc: 1.0000 - val
_loss: 0.0754 - val_acc: 0.9857
Epoch 16/20
60000/60000 [=====] - 1s 22us/step - loss: 1.1272e-04 - acc: 1.0000 - val
_loss: 0.0771 - val_acc: 0.9851
Epoch 17/20
60000/60000 [=====] - 1s 22us/step - loss: 7.6344e-05 - acc: 1.0000 - val
_loss: 0.0780 - val_acc: 0.9852
Epoch 18/20
60000/60000 [=====] - 1s 24us/step - loss: 6.4472e-05 - acc: 1.0000 - val
_loss: 0.0787 - val_acc: 0.9853
Epoch 19/20
60000/60000 [=====] - 2s 25us/step - loss: 5.5575e-05 - acc: 1.0000 - val
_loss: 0.0803 - val_acc: 0.9852
Epoch 20/20
60000/60000 [=====] - 2s 25us/step - loss: 5.0057e-05 - acc: 1.0000 - val
_loss: 0.0807 - val_acc: 0.9857

```

In [40]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
score7=score[0]
score8=score[1]
train_acc4=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax21 = plt.subplots(1,1)
ax21.set_xlabel('epoch') ; ax21.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))


vy21 = history21.history['val_loss']
ty21 = history21.history['loss']
plt_dynamic(x, vy21, ty21, ax21)

```

Test score: 0.08068556088549227

Test accuracy: 0.9857





In [41]:

```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2.2 MLP + Batch-Norm on hidden Layers + AdamOptimizer

In [42]:

```

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_relu.add(Dense(352, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())
model_relu.add(Dense(164, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_relu.add(Dense(124, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))
model_batch.compile(optimizer='adam', loss='categorical_crossentropy',
                    metrics=['accuracy'])

history22 = model_batch.fit(X_train, Y_train,
                            batch_size=batch_size,
                            epochs=nb_epoch, verbose=1,
                            validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 3s 46us/step - loss: 0.4860 - acc: 0.8565 -
val_loss: 0.6053 - val_acc: 0.8472
Epoch 2/20
60000/60000 [=====] - 2s 30us/step - loss: 0.3029 - acc: 0.9134 -
val_loss: 2.6513 - val_acc: 0.4629
Epoch 3/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2857 - acc: 0.9189 -
val_loss: 0.8559 - val_acc: 0.7858
Epoch 4/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2774 - acc: 0.9219 -
val_loss: 0.8852 - val_acc: 0.7914
Epoch 5/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2712 - acc: 0.9236 -
val_loss: 7.8901 - val_acc: 0.1965
Epoch 6/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2689 - acc: 0.9243 -
val_loss: 1.0264 - val_acc: 0.7722
Epoch 7/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2669 - acc: 0.9258 -
val_loss: 1.1807 - val_acc: 0.7158
Epoch 8/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2635 - acc: 0.9253 -
val_loss: 1.4079 - val_acc: 0.6595
Epoch 9/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2624 - acc: 0.9263 -
val_loss: 1.6847 - val_acc: 0.6442
Epoch 10/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2591 - acc: 0.9277 -
val_loss: 2.0835 - val_acc: 0.5957
Epoch 11/20
60000/60000 [=====] - 2s 30us/step - loss: 0.2597 - acc: 0.9271 -
val_loss: 0.7050 - val_acc: 0.8390
Epoch 12/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2570 - acc: 0.9278 -
val_loss: 5.7033 - val_acc: 0.4569
Epoch 13/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2560 - acc: 0.9289 -
val_loss: 2.0083 - val_acc: 0.5987
Epoch 14/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2576 - acc: 0.9276 -
val_loss: 1.5552 - val_acc: 0.6973
Epoch 15/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2544 - acc: 0.9284 -
val_loss: 9.8770 - val_acc: 0.1527
Epoch 16/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2553 - acc: 0.9274 -
val_loss: 0.7170 - val_acc: 0.8289
Epoch 17/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2518 - acc: 0.9290 -
val_loss: 3.4736 - val_acc: 0.4885
Epoch 18/20

```

```
Epoch 18/20
60000/60000 [=====] - 2s 34us/step - loss: 0.2524 - acc: 0.9293 -
val_loss: 2.1239 - val_acc: 0.5799
Epoch 19/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2517 - acc: 0.9287 -
val_loss: 0.6194 - val_acc: 0.8697
Epoch 20/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2522 - acc: 0.9295 -
val_loss: 1.0895 - val_acc: 0.7509
```

In [43]:

```
model_batch.summary()
```

Layer (type)	Output Shape	Param #
batch_normalization_9 (Batch Normalization)	(None, 784)	3136
batch_normalization_10 (Batch Normalization)	(None, 784)	3136
batch_normalization_11 (Batch Normalization)	(None, 784)	3136
dense_30 (Dense)	(None, 10)	7850

Total params: 17,258
 Trainable params: 12,554
 Non-trainable params: 4,704

In [44]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
score9=score[0]
score10=score[1]
train_acc5=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax22 = plt.subplots(1,1)
ax22.set_xlabel('epoch') ; ax22.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy22 = history22.history['val_loss']
ty22 = history22.history['loss']
plt_dynamic(x, vy22, ty22, ax22)
```

Test score: 1.089544097495079

Test accuracy: 0.7509





In [45]:

```
w_after = model_relu.get_weights()

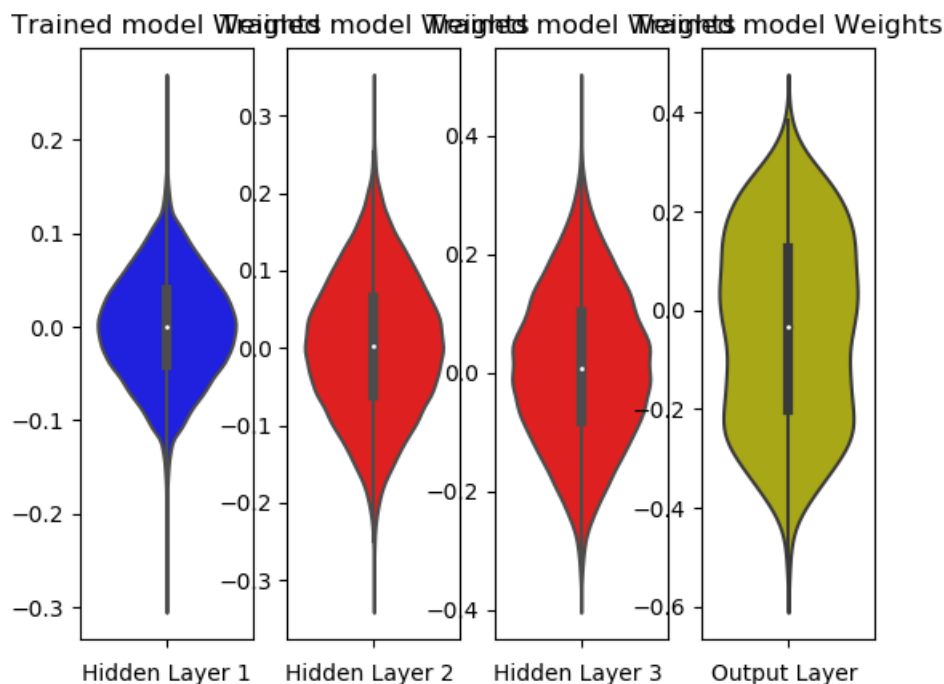
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



In [46]:

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_relu.add(Dense(352, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())
model_relu.add(Dense(164, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_relu.add(Dense(124, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))
model_batch.compile(optimizer='rmsprop', loss='categorical_crossentropy',
                    metrics=['accuracy'])

history22 = model_batch.fit(X_train, Y_train,
                            batch_size=batch_size,
                            epochs=nb_epoch, verbose=1,
                            validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 3s 46us/step - loss: 0.4528 - acc: 0.8681 -
val_loss: 0.8893 - val_acc: 0.7956
Epoch 2/20
60000/60000 [=====] - 2s 30us/step - loss: 0.3014 - acc: 0.9146 -
val_loss: 0.7591 - val_acc: 0.7998
Epoch 3/20
60000/60000 [=====] - 2s 28us/step - loss: 0.2882 - acc: 0.9201 -
val_loss: 0.7948 - val_acc: 0.7697
Epoch 4/20
60000/60000 [=====] - 2s 28us/step - loss: 0.2804 - acc: 0.9223 -
val_loss: 1.9841 - val_acc: 0.6465
Epoch 5/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2769 - acc: 0.9234 -
val_loss: 1.8222 - val_acc: 0.5741
Epoch 6/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2731 - acc: 0.9244 -
val_loss: 2.1971 - val_acc: 0.5198
Epoch 7/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2717 - acc: 0.9256 -
val_loss: 0.9766 - val_acc: 0.8075
Epoch 8/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2707 - acc: 0.9262 -
val_loss: 0.7744 - val_acc: 0.8175
Epoch 9/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2674 - acc: 0.9268 -
val_loss: 1.1930 - val_acc: 0.7625
Epoch 10/20
60000/60000 [=====] - 2s 28us/step - loss: 0.2660 - acc: 0.9273 -
val_loss: 0.6178 - val_acc: 0.8532
Epoch 11/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2664 - acc: 0.9272 -
val_loss: 2.9305 - val_acc: 0.5530
Epoch 12/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2630 - acc: 0.9278 -
val_loss: 1.0858 - val_acc: 0.7395
Epoch 13/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2635 - acc: 0.9283 -
val_loss: 0.8976 - val_acc: 0.8095
Epoch 14/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2624 - acc: 0.9295 -
val_loss: 1.7319 - val_acc: 0.6818
Epoch 15/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2624 - acc: 0.9292 -
val_loss: 3.3493 - val_acc: 0.5509
Epoch 16/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2602 - acc: 0.9294 -
val_loss: 0.6959 - val_acc: 0.8367
```

```
Epoch 17/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2591 - acc: 0.9293 -
val_loss: 1.6922 - val_acc: 0.6514
Epoch 18/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2607 - acc: 0.9294 -
val_loss: 4.0993 - val_acc: 0.4954
Epoch 19/20
60000/60000 [=====] - 2s 30us/step - loss: 0.2604 - acc: 0.9295 -
val_loss: 0.7377 - val_acc: 0.8301
Epoch 20/20
60000/60000 [=====] - 2s 30us/step - loss: 0.2577 - acc: 0.9292 -
val_loss: 2.1088 - val_acc: 0.5618
```

In [47]:


```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
score9=score[0]
score10=score[1]
train_acc5=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax22 = plt.subplots(1,1)
ax22.set_xlabel('epoch') ; ax22.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy22 = history22.history['val_loss']
ty22 = history22.history['loss']
plt_dynamic(x, vy22, ty22, ax22)
```

Test score: 2.1087711837768555
Test accuracy: 0.5618



In [48]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)
```

```

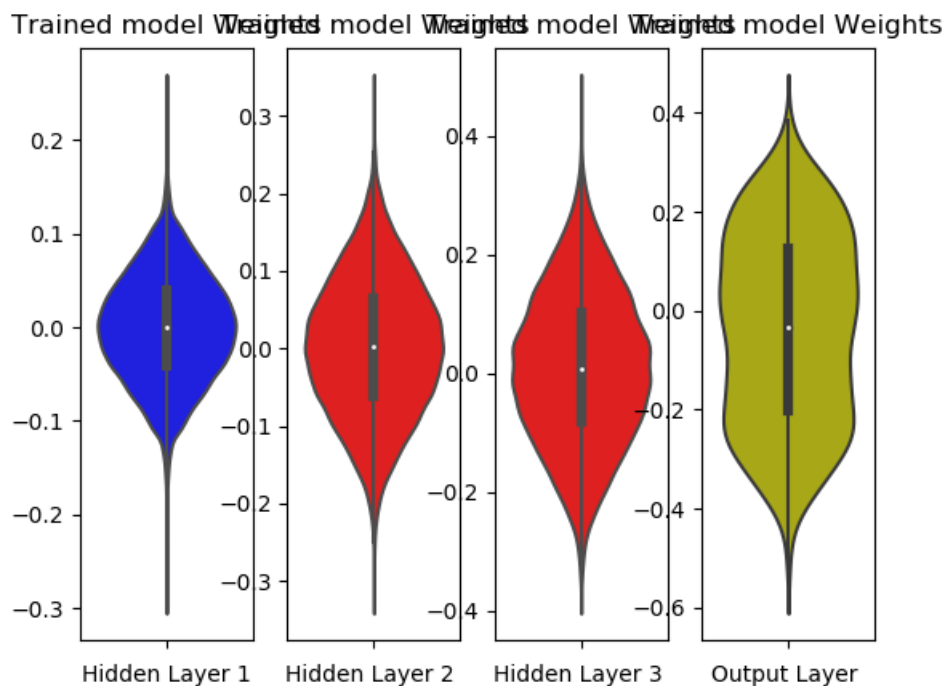
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



2.3 MLP + Dropout + AdamOptimizer

In [49]:

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()
model_relu.add(Dense(352, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))

model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_relu.add(Dense(164, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

```

```

model_relu.add(Dense(124, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

```

In [50]:

```

model_drop.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history23 = model_drop.fit(X_train, Y_train,
                          batch_size=batch_size,

                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 3s 54us/step - loss: 1.3786 - acc: 0.5656 -
val_loss: 0.5066 - val_acc: 0.8767
Epoch 2/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8568 - acc: 0.7204 -
val_loss: 0.4589 - val_acc: 0.8899
Epoch 3/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8374 - acc: 0.7253 -
val_loss: 0.4385 - val_acc: 0.8927
Epoch 4/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8287 - acc: 0.7279 -
val_loss: 0.4260 - val_acc: 0.8960
Epoch 5/20
60000/60000 [=====] - 2s 33us/step - loss: 0.8302 - acc: 0.7317 -
val_loss: 0.4264 - val_acc: 0.8952
Epoch 6/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8247 - acc: 0.7319 -
val_loss: 0.4230 - val_acc: 0.8946
Epoch 7/20
60000/60000 [=====] - 2s 35us/step - loss: 0.8203 - acc: 0.7321 -
val_loss: 0.4159 - val_acc: 0.8973
Epoch 8/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8212 - acc: 0.7347 -
val_loss: 0.4140 - val_acc: 0.8924
Epoch 9/20
60000/60000 [=====] - 2s 37us/step - loss: 0.8182 - acc: 0.7335 -
val_loss: 0.4209 - val_acc: 0.8962
Epoch 10/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8195 - acc: 0.7338 -
val_loss: 0.4100 - val_acc: 0.8968
Epoch 11/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8209 - acc: 0.7330 -
val_loss: 0.4054 - val_acc: 0.8982
Epoch 12/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8197 - acc: 0.7331 -
val_loss: 0.4080 - val_acc: 0.8968
Epoch 13/20
60000/60000 [=====] - 2s 35us/step - loss: 0.8123 - acc: 0.7372 -
val_loss: 0.4037 - val_acc: 0.8977
Epoch 14/20
60000/60000 [=====] - 2s 35us/step - loss: 0.8109 - acc: 0.7377 -
val_loss: 0.4071 - val_acc: 0.8982
Epoch 15/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8165 - acc: 0.7362 -
val_loss: 0.4074 - val_acc: 0.8971
Epoch 16/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8080 - acc: 0.7385 -
val_loss: 0.4075 - val_acc: 0.8977
Epoch 17/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8045 - acc: 0.7400 -
val_loss: 0.4034 - val_acc: 0.8983
Epoch 18/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8125 - acc: 0.7369 -
val_loss: 0.4040 - val_acc: 0.8983

```

```

val_loss: 0.4040 - val_acc: 0.8903
Epoch 19/20
60000/60000 [=====] - 2s 34us/step - loss: 0.8083 - acc: 0.7378 -
val_loss: 0.4047 - val_acc: 0.8975
Epoch 20/20
60000/60000 [=====] - 2s 35us/step - loss: 0.8052 - acc: 0.7413 -
val_loss: 0.4010 - val_acc: 0.8984

```

In [51]:

```
model_drop.summary()
```

Layer (type)	Output Shape	Param #
batch_normalization_15 (Batch Normalization)	(None, 784)	3136
dropout_5 (Dropout)	(None, 784)	0
batch_normalization_16 (Batch Normalization)	(None, 784)	3136
dropout_6 (Dropout)	(None, 784)	0
batch_normalization_17 (Batch Normalization)	(None, 784)	3136
dropout_7 (Dropout)	(None, 784)	0
dense_38 (Dense)	(None, 10)	7850
Total params: 17,258		
Trainable params: 12,554		
Non-trainable params: 4,704		

In [52]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
score1=score[0]
score12=score[1]
train_acc6=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax23 = plt.subplots(1,1)
ax23.set_xlabel('epoch') ; ax23.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy23 = history23.history['val_loss']
ty23 = history23.history['loss']
plt_dynamic(x, vy23, ty23, ax23)

```

Test score: 0.40103902480602266
Test accuracy: 0.8984





In [53]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Trained model Weights model Weights model Weights model Weights



MLP + with Different dropout rates + RMS Prop Optimizer

In [54]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout

model_drop = Sequential()
model_drop.add(Dense(352, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))

model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(164, activation='relu',
                    kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(124, activation='relu',
                    kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.4))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history23 = model_drop.fit(X_train, Y_train,
                          batch_size=batch_size,

                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 3s 53us/step - loss: 1.0449 - acc: 0.6674 -
val_loss: 0.4368 - val_acc: 0.8891
Epoch 2/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6637 - acc: 0.7915 -
val_loss: 0.3845 - val_acc: 0.8983
Epoch 3/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6323 - acc: 0.8020 -
val_loss: 0.3641 - val_acc: 0.9003
Epoch 4/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6329 - acc: 0.8011 -
val_loss: 0.3580 - val_acc: 0.9036
Epoch 5/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6290 - acc: 0.8030 -
val_loss: 0.3571 - val_acc: 0.9047
Epoch 6/20
60000/60000 [=====] - 2s 33us/step - loss: 0.6293 - acc: 0.8032 -
val_loss: 0.3534 - val_acc: 0.9050
Epoch 7/20
60000/60000 [=====] - 2s 33us/step - loss: 0.6248 - acc: 0.8047 -
val_loss: 0.3550 - val_acc: 0.9031
Epoch 8/20
60000/60000 [=====] - 2s 36us/step - loss: 0.6191 - acc: 0.8059 -
val_loss: 0.3503 - val_acc: 0.9058
Epoch 9/20
60000/60000 [=====] - 2s 33us/step - loss: 0.6176 - acc: 0.8060 -
val_loss: 0.3455 - val_acc: 0.9072
Epoch 10/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6211 - acc: 0.8079 -
val_loss: 0.3505 - val_acc: 0.9052
Epoch 11/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6208 - acc: 0.8065 -
val_loss: 0.3435 - val_acc: 0.9051
Epoch 12/20
```



```

Epoch 12/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6160 - acc: 0.8074 -
val_loss: 0.3414 - val_acc: 0.9057
Epoch 13/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6151 - acc: 0.8092 -
val_loss: 0.3399 - val_acc: 0.9087
Epoch 14/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6223 - acc: 0.8082 -
val_loss: 0.3429 - val_acc: 0.9075
Epoch 15/20
60000/60000 [=====] - 2s 33us/step - loss: 0.6183 - acc: 0.8070 -
val_loss: 0.3386 - val_acc: 0.9094
Epoch 16/20
60000/60000 [=====] - 2s 32us/step - loss: 0.6078 - acc: 0.8110 -
val_loss: 0.3395 - val_acc: 0.9069
Epoch 17/20
60000/60000 [=====] - 2s 33us/step - loss: 0.6131 - acc: 0.8093 -
val_loss: 0.3380 - val_acc: 0.9068
Epoch 18/20
60000/60000 [=====] - 2s 33us/step - loss: 0.6169 - acc: 0.8086 -
val_loss: 0.3439 - val_acc: 0.9047
Epoch 19/20
60000/60000 [=====] - 2s 33us/step - loss: 0.6122 - acc: 0.8114 -
val_loss: 0.3397 - val_acc: 0.9092
Epoch 20/20
60000/60000 [=====] - 2s 33us/step - loss: 0.6135 - acc: 0.8109 -
val_loss: 0.3385 - val_acc: 0.9068

```

In [55]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
score1=score[0]
score2=score[1]
train_acc6=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax23 = plt.subplots(1,1)
ax23.set_xlabel('epoch') ; ax23.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))


vy23 = history23.history['val_loss']
ty23 = history23.history['loss']
plt_dynamic(x, vy23, ty23, ax23)

```

Test score: 0.33852629685401914

Test accuracy: 0.9068





In [56]:

```
w_after = model_relu.get_weights()

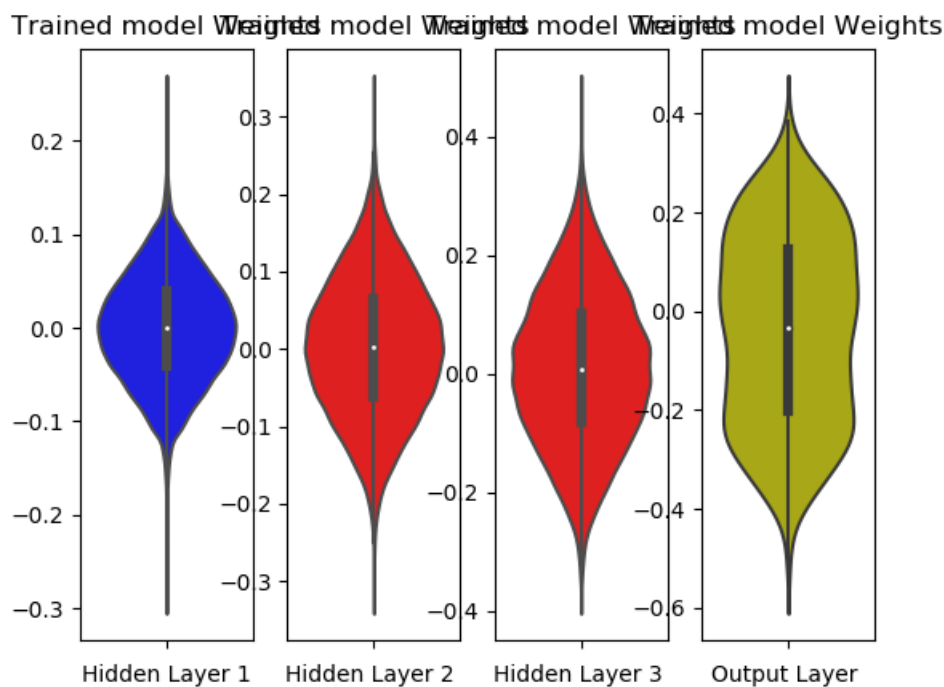
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3) 5-Hidden layer architecture (784-216-170-136-80-38-10 architecture)

3.1 MLP + ReLU + ADAM

In [57]:

```
model_relu = Sequential()
model_relu.add(Dense(216, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_relu.add(Dense(170, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )

model_relu.add(Dense(136, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(80, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )

model_relu.add(Dense(38, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history31 = model_relu.fit(X_train, Y_train,
                          batch_size=batch_size,
                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_43 (Dense)	(None, 216)	169560
dense_44 (Dense)	(None, 170)	36890
dense_45 (Dense)	(None, 136)	23256
dense_46 (Dense)	(None, 80)	10960
dense_47 (Dense)	(None, 38)	3078
dense_48 (Dense)	(None, 10)	390

Total params: 244,134
Trainable params: 244,134
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 48us/step - loss: 0.2795 - acc: 0.9164 -
val_loss: 0.1212 - val_acc: 0.9622

Epoch 2/20

60000/60000 [=====] - 2s 27us/step - loss: 0.1043 - acc: 0.9682 -
val_loss: 0.1048 - val_acc: 0.9681

Epoch 3/20

60000/60000 [=====] - 2s 26us/step - loss: 0.0724 - acc: 0.9777 -
val_loss: 0.0897 - val_acc: 0.9719

Epoch 4/20

60000/60000 [=====] - 2s 26us/step - loss: 0.0546 - acc: 0.9829 -
val_loss: 0.0902 - val_acc: 0.9736

Epoch 5/20

60000/60000 [=====] - 2s 27us/step - loss: 0.0425 - acc: 0.9867 -
val_loss: 0.0855 - val_acc: 0.9751

Epoch 6/20

60000/60000 [=====] - 2s 25us/step - loss: 0.0352 - acc: 0.9886 -
val_loss: 0.0867 - val_acc: 0.9770

Epoch 7/20

60000/60000 [=====] - 2s 25us/step - loss: 0.0343 - acc: 0.9890 -
val_loss: 0.0845 - val_acc: 0.9758

Epoch 8/20

60000/60000 [=====] - 2s 26us/step - loss: 0.0269 - acc: 0.9910 -
val_loss: 0.0978 - val_acc: 0.9738

Epoch 9/20

```

60000/60000 [=====] - 2s 26us/step - loss: 0.0247 - acc: 0.9919 -
val_loss: 0.0975 - val_acc: 0.9766
Epoch 10/20
60000/60000 [=====] - 2s 28us/step - loss: 0.0208 - acc: 0.9929 -
val_loss: 0.1001 - val_acc: 0.9743
Epoch 11/20
60000/60000 [=====] - 2s 29us/step - loss: 0.0209 - acc: 0.9932 -
val_loss: 0.0882 - val_acc: 0.9775
Epoch 12/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0185 - acc: 0.9943 -
val_loss: 0.0866 - val_acc: 0.9787
Epoch 13/20
60000/60000 [=====] - 2s 29us/step - loss: 0.0164 - acc: 0.9947 -
val_loss: 0.1214 - val_acc: 0.9752
Epoch 14/20
60000/60000 [=====] - 2s 30us/step - loss: 0.0186 - acc: 0.9940 -
val_loss: 0.0973 - val_acc: 0.9767
Epoch 15/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0147 - acc: 0.9952 -
val_loss: 0.0901 - val_acc: 0.9800
Epoch 16/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0151 - acc: 0.9953 -
val_loss: 0.0932 - val_acc: 0.9779
Epoch 17/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0149 - acc: 0.9952 -
val_loss: 0.1049 - val_acc: 0.9758
Epoch 18/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0121 - acc: 0.9960 -
val_loss: 0.1006 - val_acc: 0.9788
Epoch 19/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0119 - acc: 0.9963 -
val_loss: 0.0871 - val_acc: 0.9816
Epoch 20/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0129 - acc: 0.9961 -
val_loss: 0.0933 - val_acc: 0.9778

```

In [58]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
score13=score[0]
score14=score[1]
train_acc7=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax31 = plt.subplots(1,1)
ax31.set_xlabel('epoch') ; ax31.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy31 = history31.history['val_loss']
ty31 = history31.history['loss']
plt_dynamic(x, vy31, ty31, ax31)

```

Test score: 0.09329951099789031
Test accuracy: 0.9778





In [60]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')


plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Trained model Weights, Hidden Layer 1, Hidden Layer 2, Hidden Layer 3, Hidden Layer 4, Hidden Layer 5, Output Layer





MLP + ReLU + rmsprop

In [61]:

```
model_relu = Sequential()
model_relu.add(Dense(216, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_relu.add(Dense(170, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )

model_relu.add(Dense(136, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(80, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )

model_relu.add(Dense(38, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history31 = model_relu.fit(X_train, Y_train,
                          batch_size=batch_size,
                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_49 (Dense)	(None, 216)	169560
dense_50 (Dense)	(None, 170)	36890
dense_51 (Dense)	(None, 136)	23256
dense_52 (Dense)	(None, 80)	10960
dense_53 (Dense)	(None, 38)	3078
dense_54 (Dense)	(None, 10)	390
Total params: 244,134		
Trainable params: 244,134		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 44us/step - loss: 0.3035 - acc: 0.9089 -

val_loss: 0.1484 - val_acc: 0.9545

Epoch 2/20

60000/60000 [=====] - 2s 26us/step - loss: 0.1106 - acc: 0.9665 -

val_loss: 0.1382 - val_acc: 0.9611

Epoch 3/20

60000/60000 [=====] - 2s 25us/step - loss: 0.0790 - acc: 0.9763 -

val_loss: 0.1039 - val_acc: 0.9671

Epoch 4/20

```

Epoch 4/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0602 - acc: 0.9813 -
val_loss: 0.0856 - val_acc: 0.9762
Epoch 5/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0487 - acc: 0.9856 -
val_loss: 0.1027 - val_acc: 0.9721
Epoch 6/20
60000/60000 [=====] - 2s 25us/step - loss: 0.0387 - acc: 0.9877 -
val_loss: 0.1008 - val_acc: 0.9748
Epoch 7/20
60000/60000 [=====] - 2s 27us/step - loss: 0.0322 - acc: 0.9900 -
val_loss: 0.0860 - val_acc: 0.9792
Epoch 8/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0286 - acc: 0.9913 -
val_loss: 0.1161 - val_acc: 0.9760
Epoch 9/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0253 - acc: 0.9924 -
val_loss: 0.0966 - val_acc: 0.9795
Epoch 10/20
60000/60000 [=====] - 2s 30us/step - loss: 0.0233 - acc: 0.9934 -
val_loss: 0.1150 - val_acc: 0.9767
Epoch 11/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0207 - acc: 0.9940 -
val_loss: 0.1069 - val_acc: 0.9789
Epoch 12/20
60000/60000 [=====] - 2s 30us/step - loss: 0.0201 - acc: 0.9943 -
val_loss: 0.1299 - val_acc: 0.9768
Epoch 13/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0171 - acc: 0.9953 -
val_loss: 0.1255 - val_acc: 0.9801
Epoch 14/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0172 - acc: 0.9949 -
val_loss: 0.1009 - val_acc: 0.9813
Epoch 15/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0145 - acc: 0.9959 -
val_loss: 0.1383 - val_acc: 0.9790
Epoch 16/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0156 - acc: 0.9959 -
val_loss: 0.1256 - val_acc: 0.9781
Epoch 17/20
60000/60000 [=====] - 2s 26us/step - loss: 0.0143 - acc: 0.9964 -
val_loss: 0.1326 - val_acc: 0.9807
Epoch 18/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0120 - acc: 0.9969 -
val_loss: 0.1784 - val_acc: 0.9779
Epoch 19/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0121 - acc: 0.9968 -
val_loss: 0.1648 - val_acc: 0.9780
Epoch 20/20
60000/60000 [=====] - 1s 24us/step - loss: 0.0125 - acc: 0.9968 -
val_loss: 0.1364 - val_acc: 0.9792

```

In [62]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
score13=score[0]
score14=score[1]
train_acc7=history11.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax31 = plt.subplots(1,1)
ax31.set_xlabel('epoch') ; ax31.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy31 = history31.history['val_loss']
ty31 = history31.history['loss']
plt_dynamic(x, vy31, ty31, ax31)

```

Test score: 0.13639087163173536
Test accuracy: 0.9792



In [63]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

3.2 MLP + Batch-Norm on hidden Layers + AdamOptimizer

In [64]:

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_relu.add(Dense(216, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())
model_relu.add(Dense(170, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_relu.add(Dense(136, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())
model_relu.add(Dense(80, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_relu.add(Dense(38, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))
```

In [65]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy',
                    metrics=['accuracy'])

history32 = model_batch.fit(X_train, Y_train,
                            batch_size=batch_size,
                            epochs=nb_epoch, verbose=1,
                            validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 74us/step - loss: 0.4931 - acc: 0.8537 -
val_loss: 14.2806 - val_acc: 0.1140

Epoch 2/20

60000/60000 [=====] - 3s 47us/step - loss: 0.3048 - acc: 0.9124 -

```

val_loss: 14.5466 - val_acc: 0.0975
Epoch 3/20
60000/60000 [=====] - 3s 43us/step - loss: 0.2858 - acc: 0.9185 -
val_loss: 14.4644 - val_acc: 0.1026
Epoch 4/20
60000/60000 [=====] - 3s 46us/step - loss: 0.2785 - acc: 0.9198 -
val_loss: 14.4126 - val_acc: 0.1058
Epoch 5/20
60000/60000 [=====] - 3s 43us/step - loss: 0.2734 - acc: 0.9234 -
val_loss: 14.2581 - val_acc: 0.1154
Epoch 6/20
60000/60000 [=====] - 3s 43us/step - loss: 0.2683 - acc: 0.9232 -
val_loss: 14.2887 - val_acc: 0.1135
Epoch 7/20
60000/60000 [=====] - 3s 43us/step - loss: 0.2662 - acc: 0.9248 -
val_loss: 14.5337 - val_acc: 0.0983
Epoch 8/20
60000/60000 [=====] - 3s 43us/step - loss: 0.2645 - acc: 0.9260 -
val_loss: 14.2887 - val_acc: 0.1135
Epoch 9/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2615 - acc: 0.9261 -
val_loss: 14.5498 - val_acc: 0.0973
Epoch 10/20
60000/60000 [=====] - 3s 44us/step - loss: 0.2603 - acc: 0.9273 -
val_loss: 14.4902 - val_acc: 0.1010
Epoch 11/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2595 - acc: 0.9271 -
val_loss: 14.2855 - val_acc: 0.1137
Epoch 12/20
60000/60000 [=====] - 3s 45us/step - loss: 0.2561 - acc: 0.9278 -
val_loss: 14.4740 - val_acc: 0.1020
Epoch 13/20
60000/60000 [=====] - 3s 44us/step - loss: 0.2563 - acc: 0.9276 -
val_loss: 14.4878 - val_acc: 0.1011
Epoch 14/20
60000/60000 [=====] - 3s 46us/step - loss: 0.2550 - acc: 0.9285 -
val_loss: 14.5450 - val_acc: 0.0976
Epoch 15/20
60000/60000 [=====] - 3s 44us/step - loss: 0.2546 - acc: 0.9289 -
val_loss: 14.4660 - val_acc: 0.1025
Epoch 16/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2529 - acc: 0.9285 -
val_loss: 14.5514 - val_acc: 0.0972
Epoch 17/20
60000/60000 [=====] - 3s 51us/step - loss: 0.2541 - acc: 0.9288 -
val_loss: 14.4886 - val_acc: 0.1011
Epoch 18/20
60000/60000 [=====] - 3s 44us/step - loss: 0.2515 - acc: 0.9299 -
val_loss: 14.6723 - val_acc: 0.0897
Epoch 19/20
60000/60000 [=====] - 3s 46us/step - loss: 0.2518 - acc: 0.9300 -
val_loss: 14.2887 - val_acc: 0.1135
Epoch 20/20
60000/60000 [=====] - 3s 44us/step - loss: 0.2533 - acc: 0.9285 -
val_loss: 14.5466 - val_acc: 0.0975

```

In [66]:

```
model_batch.summary()
```

Layer (type)	Output Shape	Param #
batch_normalization_21 (Batch Normalization)	(None, 784)	3136
batch_normalization_22 (Batch Normalization)	(None, 784)	3136
batch_normalization_23 (Batch Normalization)	(None, 784)	3136
batch_normalization_24 (Batch Normalization)	(None, 784)	3136
batch_normalization_25 (Batch Normalization)	(None, 784)	3136
dense_60 (Dense)	(None, 10)	7850

Total params: 23,530
Trainable params: 15,690
Non-trainable params: 7,840

In [67]:


```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
score15=score[0]
score16=score[1]
train_acc8=history32.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax32 = plt.subplots(1,1)
ax32.set_xlabel('epoch') ; ax32.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy32 = history32.history['val_loss']
ty32 = history32.history['loss']
plt_dynamic(x, vy32, ty32, ax32)
```

Test score: 14.546580853271484
Test accuracy: 0.0975



In [68]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
```

```
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on hidden Layers + Adadelta

In [69]:

```
model_batch = Sequential()

model_relu.add(Dense(216, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_batch.add(BatchNormalization())
model_relu.add(Dense(170, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_relu.add(Dense(136, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())
model_relu.add(Dense(60, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())
```

```

model_relu.add(Dense(80, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_relu.add(Dense(38, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

```

In [70]:

```

model_batch.compile(optimizer='adadelta', loss='categorical_crossentropy',
                  metrics=['accuracy'])

history32 = model_batch.fit(X_train, Y_train,
                          batch_size=batch_size,
                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 5s 76us/step - loss: 0.4841 - acc: 0.8597 -
val_loss: 14.5256 - val_acc: 0.0988
Epoch 2/20
60000/60000 [=====] - 3s 58us/step - loss: 0.3068 - acc: 0.9132 -
val_loss: 14.4612 - val_acc: 0.1028
Epoch 3/20
60000/60000 [=====] - 3s 51us/step - loss: 0.2891 - acc: 0.9184 -
val_loss: 14.4525 - val_acc: 0.1032
Epoch 4/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2803 - acc: 0.9213 -
val_loss: 14.2806 - val_acc: 0.1140
Epoch 5/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2754 - acc: 0.9231 -
val_loss: 14.4886 - val_acc: 0.1011
Epoch 6/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2723 - acc: 0.9242 -
val_loss: 14.4628 - val_acc: 0.1027
Epoch 7/20
60000/60000 [=====] - 3s 49us/step - loss: 0.2694 - acc: 0.9256 -
val_loss: 14.6820 - val_acc: 0.0891
Epoch 8/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2693 - acc: 0.9252 -
val_loss: 14.5233 - val_acc: 0.0989
Epoch 9/20
60000/60000 [=====] - 3s 57us/step - loss: 0.2664 - acc: 0.9268 -
val_loss: 14.5422 - val_acc: 0.0977
Epoch 10/20
60000/60000 [=====] - 3s 46us/step - loss: 0.2642 - acc: 0.9260 -
val_loss: 14.6614 - val_acc: 0.0903
Epoch 11/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2625 - acc: 0.9281 -
val_loss: 14.4832 - val_acc: 0.1013
Epoch 12/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2619 - acc: 0.9269 -
val_loss: 14.2806 - val_acc: 0.1140
Epoch 13/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2598 - acc: 0.9288 -
val_loss: 14.2806 - val_acc: 0.1140
Epoch 14/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2590 - acc: 0.9287 -
val_loss: 14.5417 - val_acc: 0.0978
Epoch 15/20
60000/60000 [=====] - 3s 46us/step - loss: 0.2594 - acc: 0.9279 -
val_loss: 14.5337 - val_acc: 0.0983
Epoch 16/20
60000/60000 [=====] - 3s 46us/step - loss: 0.2565 - acc: 0.9295 -
val_loss: 14.4757 - val_acc: 0.1019
Epoch 17/20
60000/60000 [=====] - 3s 46us/step - loss: 0.2565 - acc: 0.9300 -
val_loss: 14.4943 - val_acc: 0.1007
Epoch 18/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2571 - acc: 0.9294 -
val_loss: 14.2790 - val_acc: 0.1141
Epoch 19/20

```

```
Epoch 19/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2569 - acc: 0.9288 -
val_loss: 14.6698 - val_acc: 0.0897
Epoch 20/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2567 - acc: 0.9287 -
val_loss: 14.5289 - val_acc: 0.0986
```

In [71]:


```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
score15=score[0]
score16=score[1]
train_acc8=history32.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax32 = plt.subplots(1,1)
ax32.set_xlabel('epoch') ; ax32.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy32 = history32.history['val_loss']
ty32 = history32.history['loss']
plt_dynamic(x, vy32, ty32, ax32)
```

Test score: 14.528850985717773
Test accuracy: 0.0986



In [72]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
```

```
plt.figure(figsize=(10, 10))
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3.3 MLP + Dropout + AdamOptimizer

In [73]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()
model_drop.add(Dense(216, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(170, activation='relu',
```

```

model_relu.add(Dense(128, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_relu.add(Dense(136, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_relu.add(Dense(80, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_relu.add(Dense(38, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

```

In [74]:

```

model_drop.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

history33 = model_drop.fit(X_train, Y_train,
                          batch_size=batch_size,

                          epochs=nb_epoch, verbose=1,
                          validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 5s 85us/step - loss: 2.1472 - acc: 0.3123 -
val_loss: 1.0609 - val_acc: 0.8265
Epoch 2/20
60000/60000 [=====] - 3s 49us/step - loss: 1.6095 - acc: 0.4413 -
val_loss: 0.9769 - val_acc: 0.8420
Epoch 3/20
60000/60000 [=====] - 3s 52us/step - loss: 1.5983 - acc: 0.4471 -
val_loss: 0.9571 - val_acc: 0.8477
Epoch 4/20
60000/60000 [=====] - 3s 50us/step - loss: 1.5921 - acc: 0.4479 -
val_loss: 0.9511 - val_acc: 0.8460
Epoch 5/20
60000/60000 [=====] - 3s 58us/step - loss: 1.5854 - acc: 0.4501 -
val_loss: 0.9482 - val_acc: 0.8460
Epoch 6/20
60000/60000 [=====] - 4s 59us/step - loss: 1.5794 - acc: 0.4549 -
val_loss: 0.9463 - val_acc: 0.8516
Epoch 7/20
60000/60000 [=====] - 3s 51us/step - loss: 1.5878 - acc: 0.4474 -
val_loss: 0.9550 - val_acc: 0.8501
Epoch 8/20
60000/60000 [=====] - 3s 58us/step - loss: 1.5732 - acc: 0.4545 -
val_loss: 0.9427 - val_acc: 0.8478
Epoch 9/20
60000/60000 [=====] - 3s 50us/step - loss: 1.5755 - acc: 0.4537 -
val_loss: 0.9401 - val_acc: 0.8461
Epoch 10/20
60000/60000 [=====] - 3s 54us/step - loss: 1.5732 - acc: 0.4532 -
val_loss: 0.9415 - val_acc: 0.8513
Epoch 11/20
60000/60000 [=====] - 3s 51us/step - loss: 1.5702 - acc: 0.4559 -
val_loss: 0.9337 - val_acc: 0.8507
Epoch 12/20
60000/60000 [=====] - 3s 50us/step - loss: 1.5701 - acc: 0.4573 -
val_loss: 0.9341 - val_acc: 0.8509
Epoch 13/20
60000/60000 [=====] - 3s 50us/step - loss: 1.5623 - acc: 0.4591 -
val_loss: 0.9330 - val_acc: 0.8524
Epoch 14/20
60000/60000 [=====] - 3s 50us/step - loss: 1.5663 - acc: 0.4601 -
val_loss: 0.9318 - val_acc: 0.8501

```



```

Epoch 15/20
60000/60000 [=====] - 3s 50us/step - loss: 1.5658 - acc: 0.4596 -
val_loss: 0.9342 - val_acc: 0.8515
Epoch 16/20
60000/60000 [=====] - 3s 56us/step - loss: 1.5580 - acc: 0.4600 -
val_loss: 0.9315 - val_acc: 0.8568
Epoch 17/20
60000/60000 [=====] - 3s 50us/step - loss: 1.5593 - acc: 0.4587 -
val_loss: 0.9311 - val_acc: 0.8504
Epoch 18/20
60000/60000 [=====] - 3s 52us/step - loss: 1.5611 - acc: 0.4602 -
val_loss: 0.9332 - val_acc: 0.8496
Epoch 19/20
60000/60000 [=====] - 3s 54us/step - loss: 1.5560 - acc: 0.4608 -
val_loss: 0.9331 - val_acc: 0.8521
Epoch 20/20
60000/60000 [=====] - 3s 50us/step - loss: 1.5554 - acc: 0.4596 -
val_loss: 0.9358 - val_acc: 0.8527

```

In [75]:

```
model_drop.summary()
```

Layer (type)	Output Shape	Param #
batch_normalization_31 (Batch Normalization)	(None, 784)	3136
dropout_11 (Dropout)	(None, 784)	0
batch_normalization_32 (Batch Normalization)	(None, 784)	3136
dropout_12 (Dropout)	(None, 784)	0
batch_normalization_33 (Batch Normalization)	(None, 784)	3136
dropout_13 (Dropout)	(None, 784)	0
batch_normalization_34 (Batch Normalization)	(None, 784)	3136
dropout_14 (Dropout)	(None, 784)	0
batch_normalization_35 (Batch Normalization)	(None, 784)	3136
dropout_15 (Dropout)	(None, 784)	0
dense_72 (Dense)	(None, 10)	7850

Total params: 23,530
 Trainable params: 15,690
 Non-trainable params: 7,840

In [76]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
score17=score[0]
score18=score[1]
train_acc9=history33.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax33 = plt.subplots(1,1)
ax33.set_xlabel('epoch') ; ax33.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy33 = history33.history['val_loss']
ty33 = history33.history['loss']
plt_dynamic(x, vy33, ty33, ax33)

```

Test score: 0.9357627007484436

Test score: 0.9997027007404400
Test accuracy: 0.8527



In [77]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
```

```
plt.show()
```



MLP + different drop out rates + rmsprop optimizer

In [78]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout

model_drop = Sequential()
model_drop.add(Dense(216, activation='relu', input_shape=(input_dim,),
                    kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(170, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(136, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
model_drop.add(Dense(80, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(38, activation='relu',
                    kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.4))

model_drop.add(Dense(output_dim, activation='softmax'))
```

In [79]:

```
model_drop.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```

```

history33 = model_drop.fit(x_train, y_train,
                           batch_size=batch_size,

                           epochs=nb_epoch, verbose=1,
                           validation_data=(X_test, Y_test))

```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 82us/step - loss: 1.8151 - acc: 0.4056 -
val_loss: 0.7996 - val_acc: 0.8484
Epoch 2/20
60000/60000 [=====] - 3s 46us/step - loss: 1.3204 - acc: 0.5537 -
val_loss: 0.7085 - val_acc: 0.8581
Epoch 3/20
60000/60000 [=====] - 3s 46us/step - loss: 1.3041 - acc: 0.5607 -
val_loss: 0.6967 - val_acc: 0.8635
Epoch 4/20
60000/60000 [=====] - 3s 46us/step - loss: 1.2945 - acc: 0.5616 -
val_loss: 0.6833 - val_acc: 0.8666
Epoch 5/20
60000/60000 [=====] - 3s 47us/step - loss: 1.2916 - acc: 0.5623 -
val_loss: 0.6804 - val_acc: 0.8624
Epoch 6/20
60000/60000 [=====] - 3s 49us/step - loss: 1.2892 - acc: 0.5648 -
val_loss: 0.6773 - val_acc: 0.8664
Epoch 7/20
60000/60000 [=====] - 3s 48us/step - loss: 1.2872 - acc: 0.5626 -
val_loss: 0.6751 - val_acc: 0.8645
Epoch 8/20
60000/60000 [=====] - 3s 47us/step - loss: 1.2836 - acc: 0.5665 -
val_loss: 0.6713 - val_acc: 0.8674
Epoch 9/20
60000/60000 [=====] - 3s 47us/step - loss: 1.2811 - acc: 0.5667 -
val_loss: 0.6704 - val_acc: 0.8646
Epoch 10/20
60000/60000 [=====] - 3s 53us/step - loss: 1.2811 - acc: 0.5662 -
val_loss: 0.6653 - val_acc: 0.8684
Epoch 11/20
60000/60000 [=====] - 4s 59us/step - loss: 1.2789 - acc: 0.5666 -
val_loss: 0.6669 - val_acc: 0.8662
Epoch 12/20
60000/60000 [=====] - 3s 56us/step - loss: 1.2793 - acc: 0.5675 -
val_loss: 0.6659 - val_acc: 0.8765
Epoch 13/20
60000/60000 [=====] - 3s 46us/step - loss: 1.2857 - acc: 0.5632 -
val_loss: 0.6651 - val_acc: 0.8716
Epoch 14/20
60000/60000 [=====] - 3s 46us/step - loss: 1.2858 - acc: 0.5653 -
val_loss: 0.6660 - val_acc: 0.8682
Epoch 15/20
60000/60000 [=====] - 3s 54us/step - loss: 1.2783 - acc: 0.5683 -
val_loss: 0.6657 - val_acc: 0.8697
Epoch 16/20
60000/60000 [=====] - 3s 51us/step - loss: 1.2762 - acc: 0.5671 -
val_loss: 0.6672 - val_acc: 0.8665
Epoch 17/20
60000/60000 [=====] - 3s 47us/step - loss: 1.2660 - acc: 0.5719 -
val_loss: 0.6576 - val_acc: 0.8665
Epoch 18/20
60000/60000 [=====] - 4s 64us/step - loss: 1.2814 - acc: 0.5652 -
val_loss: 0.6613 - val_acc: 0.8648
Epoch 19/20
60000/60000 [=====] - 3s 50us/step - loss: 1.2698 - acc: 0.5697 -
val_loss: 0.6561 - val_acc: 0.8738
Epoch 20/20
60000/60000 [=====] - 3s 48us/step - loss: 1.2824 - acc: 0.5673 -
val_loss: 0.6649 - val_acc: 0.8705

```

In [80]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
score17=score[0]
score18=score[1]
train_acc9=history33.history['acc']
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

plt.figure(figsize=(10, 10))


fig, ax33 = plt.subplots(1, 1)
ax33.set_xlabel('epoch') ; ax33.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy33 = history33.history['val_loss']
ty33 = history33.history['loss']
plt_dynamic(x, vy33, ty33, ax33)

```

Test score: 0.6648854620933533
Test accuracy: 0.8705



In [81]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')


plt.subplot(1, 6, 4)

```

```
plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Summarizing all the models performance using Pretty Table

In [82]:

```
# Please compare all your models using Prettytable library
# Please compare all your models using Prettytable library
# http://zetcode.com/python/prettytable/

from prettytable import PrettyTable

#If you get a ModuleNotFoundError error , install prettytable using: pip3 install prettytable

x = PrettyTable()
x.field_names = ["Model", "Test-accuracy"]

x.add_row(["2-Hidden layer Architecture (784-472-168-10):MLP + ReLU activation function + ADAM optimizer", 0.979])
x.add_row(["2-Hidden layer Architecture (784-472-168-10):MLP + ReLU activation function + RMSprop optimizer", 0.9809])
x.add_row(["2-Hidden layer Architecture (784-472-168-10):MLP + Batch-Norm on hidden Layers + Adam Optimizer", 0.9797])
x.add_row(["2-Hidden layer Architecture (784-472-168-10):MLP + Batch-Norm on hidden Layers + Adagrad Optimizer", 0.9843])
x.add_row(["2-Hidden layer Architecture (784-472-168-10):MLP + Dropout + AdamOptimizer", 0.9836])
x.add_row(["2-Hidden layer Architecture (784-472-168-10):MLP + Dropout + Adadelata Optimizer", 0.984])
x.add_row(["3-Hidden layer architecture (784-352-164-124 architecture):MLP+ReLU+Adam", 0.9798])
x.add_row(["3-Hidden layer architecture (784-352-164-124 architecture):MLP + ReLU + Adadelata", 0.9857])
```

```

x.add_row(["3-Hidden layer architecture (784-352-164-124 architecture):MLP + Batch-Norm on hidden
Layers + AdamOptimizer", 0.7509])
x.add_row(["3-Hidden layer architecture (784-352-164-124 architecture):MLP + Batch-Norm on hidden
Layers + RMS Prop Optimizer", 0.56])
x.add_row(["3-Hidden layer architecture (784-352-164-124 architecture):MLP + Dropout +
AdamOptimizer", 0.8984])
x.add_row(["3-Hidden layer architecture (784-352-164-124 architecture):MLP + with Different dropou
t rates + RMS Prop Optimizer", 0.9068])
x.add_row(["5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + ReLU + ADAM"
, 0.9778])
x.add_row(["5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + ReLU + rmspr
op", 0.9792])
x.add_row(["5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + Batch-Norm o
n hidden Layers + AdamOptimizer", 0.975])
x.add_row(["5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + Batch-Norm o
n hidden Layers + Adadelata", 0.986])
x.add_row(["5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + Dropout + Ad
amOptimizer", 0.8527])
x.add_row(["5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + different dr
op out rates + rmsprop optimizer", 0.8705])

print(x)

```

```

+-----+-----+
+-----+-----+
|                                     Model                                     |
| Test-acccuracy |
+-----+-----+
+-----+-----+
| 2-Hidden layer Architecture (784-472-168-10):MLP + ReLU activation function + ADAM c |
| ptimizer | 0.979 | |
| 2-Hidden layer Architecture (784-472-168-10):MLP + ReLU activation function + RMSprop |
| optimizer | 0.9809 | |
| 2-Hidden layer Architecture (784-472-168-10):MLP + Batch-Norm on hidden Layers + Adam |
| Optimizer | 0.9797 | |
| 2-Hidden layer Architecture (784-472-168-10):MLP + Batch-Norm on hidden Layers + Adagra |
| d Optimizer | 0.9843 | |
| 0.9836 | 2-Hidden layer Architecture (784-472-168-10):MLP + Dropout + AdamOptimizer |
| | 0.984 | |
| 2-Hidden layer Architecture (784-472-168-10):MLP + Dropout + Adadelata |
| Optimizer | 0.984 | |
| 3-Hidden layer architecture (784-352-164-124 architecture):MLP+ReLU+Adam |
| 0.9798 | |
| 3-Hidden layer architecture (784-352-164-124 architecture):MLP + ReLU + |
| Adadelata | 0.9857 | |
| 3-Hidden layer architecture (784-352-164-124 architecture):MLP + Batch-Norm on hidden |
| Layers + AdamOptimizer | 0.7509 | |
| 3-Hidden layer architecture (784-352-164-124 architecture):MLP + Batch-Norm on hidden |
| Layers + RMS Prop Optimizer | 0.56 | |
| 3-Hidden layer architecture (784-352-164-124 architecture):MLP + Dropout + AdamOpt |
| imizer | 0.8984 | |
| 3-Hidden layer architecture (784-352-164-124 architecture):MLP + with Different dropout |
| rates + RMS Prop Optimizer | 0.9068 | |
| 5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + ReLU + |
| ADAM | 0.9778 | |
| 5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + ReLU + r |
| msprop | 0.9792 | |
| 5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + Batch-Norm on |
| hidden Layers + AdamOptimizer | 0.975 | |
| 5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + Batch-Norm on |
| hidden Layers + Adadelata | 0.986 | |
| 5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + Dropout + Ada |
| mOptimizer | 0.8527 | |
| 5-Hidden layer architecture (784-216-170-136-80-38-10 architecture):MLP + different drop |
| out rates + rmsprop optimizer | 0.8705 | |
+-----+-----+
+-----+-----+

```

Conclusion

-- We find that the model with 3-Hidden layer architecture (784-352-164-124 architecture):MLP + ReLU + Adadelata has the highest test accuracy and it has outperformed compared to other models with different architectures.

